

Column Property Annotation using Large Language Models

Keti Korini¹[0000–0002–2158–0070] and Christian Bizer¹[0000–0003–2367–0237]

Data and Web Science Group, University of Mannheim, Mannheim, Germany
{kkorini,christian.bizer}@uni-mannheim.de

Abstract. Column property annotation (CPA), also known as column relationship prediction, is the task of predicting the semantic relationship between two columns in a table given a set of candidate relationships. CPA annotations are used in downstream tasks such as data search, data integration, or knowledge graph enrichment. This paper explores the usage of generative large language models (LLMs) for the CPA task. We experiment with different zero-shot prompts for the CPA task which we evaluate using GPT-3.5, GPT-4, and the open-source model SOLAR. We find GPT-3.5 to be quite sensitive to variations of the prompt, while GPT-4 reaches a high performance independent of the variation of the prompt. We further explore the scenario where training data for the CPA task is available and can be used for selecting demonstrations or fine-tuning the model. We show that a fine-tuned GPT-3.5 model outperforms a RoBERTa model that was fine-tuned on the same data by 11% in F1. Comparing in-context learning via demonstrations and fine-tuning shows that the fine-tuned GPT-3.5 performs 9% F1 better than the same model given demonstrations. The fine-tuned GPT-3.5 model also outperforms zero-shot GPT-4 by around 2% F1 for the dataset on which it was fine-tuned, while not generalizing to tasks that require a different vocabulary.

Keywords: Table Annotation · Large Language Models · Column Property Annotation

1 Introduction

Table annotation is the task of annotating elements of a table using terms from a pre-defined vocabulary in order to discover their semantics [14]. It consists of several sub-tasks that aim at discovering the semantics of different elements of the table. Two of the sub-tasks are column property annotation (CPA) which focuses on discovering the semantic relationship between two columns, and column type annotation (CTA) which aims at discovering the semantic type of entities contained in a column. Figure 1 shows an example of both tasks. The example table describes books, the names of which are contained in the first column and some of their attributes are contained in the other three columns. The goal of a CTA system is to discover the types of each column separately, for example the last column contains dates therefore the CTA label assigned to this

column by the system would be *Date*. The goal of a CPA system is to annotate the relationships of the columns with the first column of the table, also referred to as the subject column [21]. As the last column contains the dates when the books listed in the first column were published, a CPA system would annotate this relationship with the label *datePublished*.

BookName	BookFormatType	Language	Date
A Handbook for Morning Time	Paperback	English	01-01-2016
The Intentional Brain	Hardcover	English	15-06-2016
The Comeback	Hardcover	English	03-08-2020

inLanguage

datePublished

Fig. 1. Example of CTA and CPA annotations. CTA labels are shown above the table columns, while CPA labels are shown below the table.

Early statistical approaches [12, 21] use a maximum likelihood or maximize joint probabilities to assign CTA and CPA labels to columns or pair of columns. Later systems [7, 15, 16] use a knowledge base (KB) such as DBpedia [1] to first match the entities in columns to entities in the KB and consider the KB class of the entity as its column type while considering the KB properties of the entities as potential CPA labels. Cannaviccio et al. [2] use a combination of language modeling and using a KB. Recent approaches often rely on pre-trained language model (PLM) such as BERT [4]. The approaches fall into two groups: methods that learn tabular embeddings such as TURL [3] where the authors propose an architecture that learns cell representations which can be used for predicting CTA and CPA labels, or works that fine-tune PLMs such as DODUO [19] which experiments with table serialization and a multi-task learning architecture for CTA and CPA. With the advancements in generative large language models (LLMs) [17, 20, 24], such as GPT-3.5, GPT-4, LLaMa, Gemini, and Mixtral, research has started to explore prompt designs for table tasks as well as fine-tuning these models on table tasks: [9] compares prompt designs for the CTA task. In Chorus [8], different table tasks are explored, including CTA. The authors test adding instructions to their prompts and introduce the concept of *anchoring* for mapping the answers of the model to the original label space. ArcheType [5] fine-tunes a LLaMa-7B model on the CTA task and compares its results to two PLM baselines. Table-GPT [11] fine-tunes the *text-davinci-002* GPT-3.5 model using a combination of unsupervised table tasks such as row/column filtering, row/column sorting as well as supervised table tasks such as schema matching and entity matching. They show that their fine-tuned model generalizes to other unseen tasks. In our work, we also fine-tune a GPT-3.5 model, *gpt-3.5-0613*, on the CTA and CPA task separately as well as their combination to test if fine-

tuning on both tasks provides a better generalization ability than fine-tuning on the tasks separately. In TableLlama [23] a LLaMa-7B model is fine-tuned on six table tasks, two of which are CTA and CPA. [23] does not experiment with zero-shot prompts for the CTA task nor does it fine-tune the LLM exclusively for CPA. This paper fills these gaps. The contributions of the paper are:

1. We are the first to compose prompt designs for the CPA task in a zero-shot and few-shot setting while existing work only explores fine-tuning LLMs for the CPA task amongst other tasks.
2. Using different zero- and few-shot prompts, we analyze the performance and prompt sensitivity of GPT-3.5, GPT-4, and SOLAR-70B for the CPA task.
3. Existing research has only fine-tuned for the CPA task as one task amongst others in a multi-task learning setting. In contrast, we explore the effect of fine-tuning *gpt-3.5-turbo-0613* exclusively for the CPA task and explore how the fine-tuned model generalizes to other datasets and to the CTA task.

2 Experimental Setup

This section introduces our experimental setup. We make the code and data available on GitHub¹ so that all our experiments can be replicated.

Datasets. We use two datasets for the experiments on the CPA task. The first dataset is SOTAB V2 CPA [10] which consists of tables whose topics range across 17 domains, including books, products, local businesses etc. Its test set consists of 595 tables with 2,340 columns annotated using 108 schema.org² terms which are manually verified. The second dataset is the T2Dv2 CPA dataset. The dataset was originally published by Ritze et al. [18] and we use the manually verified version³, the test set of which consists of 80 tables from domains such as animals, book, country etc. labeled using 48 terms from DBpedia.

Additionally, for the fine-tuning experiments we use two more datasets for the evaluation of the CTA task. SOTAB V2 CTA [10] consists of tables with topics ranging over 17 domains including movies, music albums, events etc. Its test set consists of 609 tables where 1851 columns are labeled using 82 schema.org terms and the annotation is manually verified. Lastly, we build the T2Dv2 CTA dataset by using T2Dv2 CPA’s tables where we map the DBpedia properties to DBpedia classes to generate the CTA labels. Statistics about all datasets can be found in Table 1. For training in our experiments, we do not use the original large training sets for SOTAB V2 CTA and CPA, but we use down-sampled train sets to explore the scenario where less training data is available.

Language models. The LLMs that we test in the zero and few-shot scenario are two of OpenAI’s GPT models⁴, *gpt-3.5-turbo-0125* and *gpt-4-0125-preview*

¹ <https://github.com/wbgs-uni-mannheim/TabAnnGPT>

² <https://schema.org/>

³ <https://webdatacommons.org/structureddata/smb/>

⁴ <https://platform.openai.com/docs/models>

Table 1. Statistics of datasets used.

Dataset	Original Train		Sampled Train		Test		Labels
	Tables	Columns	Tables	Columns	Tables	Columns	
SOTAB V2 CTA	44,769	116,887	1199	1640	609	1,851	82
T2Dv2 CTA	74	146	-	-	71	145	16
SOTAB V2 CPA	29,158	109,994	1264	2160	565	2,340	108
T2Dv2 CPA	81	170	-	-	82	166	48

and one open-source model SOLAR-70B⁵, which is a LLaMa-2-70B [20] fine-tuned model. In our experiments, we will refer to these models as GPT-3.5, GPT-4 and SOLAR respectively. For our fine-tuning experiments, we fine-tune *gpt-3.5-turbo-0613*. To build our prompt templates and to access OpenAI’s models we use the *Langchain*⁶ library, while for using the open-source model we use the *Huggingface transformers*⁷ library. In order to make our experiments reproducible, we set the temperature of the models to 0. We use one NVIDIA A100 provided by *bwhPC*⁸ to run the SOLAR experiments.

Evaluation Setup. We use a multi-class classification setup and report Micro-F1 as evaluation metric due to the imbalance of the different classes. We consider answers that do not directly mention terms from the label set as errors and do not try to map such OOV (out of vocabulary) answers to the label set.

3 Comparison of Zero-shot Prompts

This section compares the performance of different zero-shot prompts for the CPA task. We distinguish three main parts in the prompts: task description, instructions, and classification sentence. The task description part aims at describing the CPA task to the model. In the instructions part we aim at writing some simple instructions that can help the model follow the CPA tasks’ steps and inform the model of a preferred format for generating its answer. In the last part, we test how classification words can influence the answer of the model. Two example prompts are shown in Figure 2. The prompt on the left contains as its first message a formulation of the task part where we only *describe* the CPA task without mentioning the name of the task. It is followed by a five-step *instructions* part which informs the model that the input is a table and the answer should be returned in a required format. In the classification message we ask the model to classify the table columns and pass the first five rows of a table in a markdown format. If in the first five rows some values are missing, we fill the cells using values from the rest of the rows. The prompt on the right contains in the first message a task formulation where the *cpa* task is mentioned and explained. It is followed by a second message that contains *less instructions*

⁵ <https://huggingface.co/upstage/SOLAR-0-70b-16bit>

⁶ <https://www.langchain.com/>

⁷ <https://github.com/huggingface/transformers>

⁸ <https://www.bwhpc.de/cluster.php>

than the prompt on the left, where we have removed the first two steps. Finally in its last message we test the keyword *annotate* to ask the model to return the labels for the CPA task. For the last message, we also test the words *determine* and classification of *relationships*.

<p>describe</p> <p>System message: Classify the relationship between two columns of a given table with one of the following relationships that are separated with comma: publisher, genre, date published, pages, job title, release date, actor, ...</p>	<p>cpa</p> <p>System message: Your task is to perform column property annotation (CPA), meaning that your task is to predict the semantic relation between the subject column and the other columns of a given table with only one of the following relationships that are separated with comma: publisher, genre, date published, pages, job title, release date, actor ...</p>
<p>instructions</p> <p>System message: 1. Look at the input given to you and make a table out of it. 2. Look at the cell values in detail. 3. For each column, select a relationship that best represents the relationship between that column and the first column of the table. 4. Answer with only one selected relationship for each column with the format Column 2: relationship. Don't return any relationship for the first column! 5. Answer only with labels from the provided label set!</p>	<p>less-instructions</p> <p>System message: 1. For each column, select a relationship that best represents the relationship between that column and the first column of the table. 2. Answer with only one selected relationship for each column with the format Column 2: relationship. Don't return any relationship for the first column! 3. Answer only with labels from the provided label set!</p>
<p>classify</p> <p>User message: Classify these table columns: Column 1 Column 2 Column 3 Column 4 ----- ----- ----- ----- A Handbook for Morning Time Paperback English 0986325759 </p>	<p>annotate</p> <p>User message: Please annotate the columns of the following table: Column 1 Column 2 Column 3 Column 4 ----- ----- ----- ----- A Handbook for Morning Time Paperback English 0986325759 </p>

Fig. 2. Example of two CPA prompts showcasing the prompt building blocks.

Results. Table 2 reports the results of prompting the models with each combination of the three building blocks described in the above paragraph. In the cases where *cpa* is not mentioned the *describe* task description is used. The results show that including the definition of the CPA task into the prompt does not help the GPT-3.5 model. This can be seen from the combinations of the *cpa* prompt which are all below the score of 61% for both datasets. In opposite, SOLAR seems to benefit from including the task name and for both datasets the highest score is achieved in this case. We can observe, mostly from the results of T2Dv2, that simply changing the wording in the classification message improves the Micro-F1 score by up to 13% when comparing the *classify* and the *relationships* keyword. In SOTAB V2, this gap ranges from 2-4% for GPT-3.5. Overall, the results show a prompt sensitivity of 3.95 and 7.49 in the case of GPT-3.5, while when looking at the results of SOLAR and GPT-4, we notice a lower sensitivity to the different formulations, especially for GPT-4 which is the most stable model amongst the three. We calculate prompt sensitivity as the standard deviation of the scores of the different prompts. Regarding SOTAB V2, the average OOV answers for GPT-3.5 is 300, 23 for GPT-4 and for SOLAR this average is 350. For T2Dv2 the averages are 40, 1 and 5 for GPT-3.5, GPT-4 and SOLAR respectively. The token length of the “less-instr-relationships” prompt is 428,528, which results in an API usage fee of \$ 0.00012 per annotation for GPT-3.5 and \$0.0025 per annotation for GPT-4 with prices as of March 2024.

Table 2. Micro-F1 performance of different zero-shot prompts for the CPA task.

Prompt	SOTAB V2 CPA			T2Dv2 CPA		
	GPT-3.5	GPT-4	SOLAR	GPT-3.5	GPT-4	SOLAR
instr-classify	64.99	80.31	49.56	56.64	80.36	75.23
cpa-instr-classify	60.14	80.38	53.59	52.55	79.76	69.72
instr-annotate	66.06	79.88	50.56	65.08	80.00	77.30
instr-determine	64.76	80.19	50.25	64.43	82.35	77.68
instr-relationships	66.42	79.30	50.58	73.12	81.48	72.56
less-instr-classify	62.86	80.19	52.57	60.50	78.55	76.83
less-instr-annotate	65.67	79.78	52.61	62.33	81.57	75.38
less-instr-determine	64.82	79.51	52.02	56.34	81.48	73.94
less-instr-relationships	66.49	80.07	52.13	73.58	79.76	71.60
cpa-less-instr-classify	56.73	81.36	53.60	50.56	80.86	78.29
Prompt sensitivity	3.95	0.64	1.30	7.49	1.06	2.60

Table 3. Few-shot performance (Micro-F1) using two demonstration selection methods and their prompt token lengths (Tokens).

Method	shots	Tokens	SOTAB V2 CPA			T2Dv2 CPA		
			GPT-3.5	GPT-4	SOLAR	GPT-3.5	GPT-4	SOLAR
random	1	547,493	68.99	81.24	53.42	76.69	84.24	78.55
random	5	994,400	70.23	82.45	-	78.55	83.38	-
similar	1	617,256	70.68	82.62	58.21	76.13	84.77	82.67
similar	5	1,337,839	74.39	84.10	-	76.82	86.09	-

4 In-context Learning via Demonstrations

The related work explores different methods of selecting demonstrations for in-context learning [13, 22]. To conduct few-shot experiments for CPA, we employ two methods for selecting demonstrations from the training sets listed in Table 1: randomly and based on similarity. The first method randomly selects examples from the training set as demonstrations. In the similarity method, we embed the test examples without labels as well as the available training examples using the embedding model *text-embedding-ada-002*⁹ and select for each test example a number of most similar examples determined using cosine similarity. The prompt that we use for running the few-shot experiments is the *less-instr-relationships* prompt which overall gave a decent score for all models. To perform few-shot, we pass the demonstrations with the help of a message with the role of user that contains the demonstration table and an assistant role message that contains the classification output for the demonstration. These messages are passed to the model before passing the last user message which contains the test table. The results of the few-shot experiments are listed in Table 3. Due to memory issues, we could run SOLAR only in a one-shot setting. From the table, we can observe that in the case of GPT-3.5 and GPT-4 the similarity method outperforms the

⁹ <https://platform.openai.com/docs/models/embeddings>

random method by at least 2% for SOTAB V2, while for T2Dv2 the results are close for both methods. Lastly, SOLAR benefits from the similar demonstrations by 5 and 4% for SOTAB V2 and T2Dv2 respectively.

5 Fine-tuning the LLM

We fine-tune the *gpt-3.5-turbo-0613* model using the down-sampled training sets of SOTAB V2 CTA and CPA and the *instr-classify* prompt. We fine-tune the model with different combinations of training sets: First, we fine-tune the model using only the SOTAB V2 CTA training set. Second, we fine-tune the model only on the CPA dataset therefore only on the CPA task. As third approach, we fine-tune the model using the CTA and CPA datasets from the two previous steps merged together to fine-tune on both tasks simultaneously. As the last approach, we fine-tune again on both tasks but with both datasets halved to make the number of fine-tuning examples comparable with the first two setups. We refer to these models as *cta-ft*, *cpa-ft*, *cta-cpa-ft* and *cta-cpa-ft-small* respectively.

Table 4. Micro-F1 fine-tuning results on CTA and CPA compared to GPT-3.5 zero-shot results. F-Tokens reports the number of tokens of the fine-tuning dataset, while F-Cost reports the cost of fine-tuning the model with prices as of March 2024.

	F-Tokens	F-Cost	SOTAB V2		T2Dv2	
			CTA	CPA	CTA	CPA
zero-shot	-	-	64.35	66.49	69.85	73.58
cta-ft	1,787,364	\$14.92	81.22	63.98	69.43	72.34
cpa-ft	2,341,017	\$18.72	70.95	83.55	65.74	72.95
cta-cpa-ft	4,128,381	\$33	82.01	84.08	72.92	76.36
cta-cpa-ft-small	2,081,448	\$16.65	80.35	80.80	71.97	69.51

Results. The results of fine-tuning are shown in Table 4. From them we conclude that fine-tuning for a specific task significantly increases the Micro-F1 score for the task and dataset that was used for fine-tuning, while decreasing the performance on the unseen T2Dv2 datasets. When fine-tuning for both tasks with the larger set, we observe that the model generalizes better to all datasets. On the other hand, when using the small CTA and CPA dataset for fine-tuning, the resulting model performs good on only the datasets that have been used for fine-tuning and not on the unseen T2Dv2 datasets. By comparing the results in Table 3 and Table 4, we can conclude that when using GPT-3.5 it is more beneficial to use the training set for fine-tuning the model, which results in 83.55% Micro-F1, rather than using the training set as a pool for selecting demonstrations which reaches 74.39% for SOTAB V2. In addition, fine-tuning also helps reduce the number of OOV answers to around 20-30 for SOTAB V2.

6 Comparison to PLM Baselines

We compare the previous zero-shot prompts for GPT-3.5, GPT-4 and fine-tuned GPT-3.5 (FT-GPT-3.5) to three baselines. The first baseline is a fine-tuned RoBERTa model with the maximum token length set to 512 and a batch size of 32. The pairs of columns are concatenated together and passed to RoBERTa which we train for 30 epochs. The second baseline TURL [3] is pre-trained using table corpora and uses TinyBERT [6] in its architecture. For this baseline we use CrossEntropy as a loss function and train the model for 50 epochs. The last PLM baseline DODUO [19] uses a new serialization method where a single table is serialized into one sequence. For this model, we use a batch size of 32 and run training for 30 epochs. We train all the baselines using a learning rate of 5e-5 and report the average performance of three runs with different random seeds.

Table 5. F1 results of fine-tuned PLM baselines compared to zero-shot results of LLMs.

Method	shots	SOTAB V2 CPA	shots	T2Dv2 CPA
GPT-3.5	0	66.49	0	73.58
GPT-4	0	81.43	0	82.35
FT-GPT-3.5	2160	83.55	2160	72.95
TURL	2160	60.75	170	59.23
DODUO	10,496	70.34	170	4.08
RoBERTa	2160	71.45	170	81.52

Results. The results of the PLM baselines are summarized in Table 5. Comparing RoBERTa and FT-GPT-3.5 which were both fine-tuned on the SOTAB V2 CPA dataset, GPT-3.5 achieves 11% higher in Micro-F1 than RoBERTa. On the other hand, for the T2Dv2 dataset, FT-GPT-3.5 fine-tuned on SOTAB V2 does not generalize well and the score compared to RoBERTa fine-tuned with the full T2Dv2 training set is 8% lower. Comparing to TURL, FT-GPT-3.5 achieves 23% higher Micro-F1, while comparing it to DODUO which is fine-tuned with more data, FT-GPT-3.5 still achieves 11% more in F1.

7 Conclusion

Our experiments using different zero-shot prompts have shown that GPT-4 reaches a F1 score of around 80% for the CPA task, outperforming the other models while also being less sensitive to variations of the prompt. Fine-tuning GPT-3.5 for the CPA task enables the model to reach a similar performance as GPT-4 while costing less API usage fees per annotation. Our experiments further show that if training data is available, it is better to use the data for fine-tuning rather than as a pool for choosing demonstrations.

Acknowledgement. The authors acknowledge support by the state of Baden-Württemberg through bwHPC for running the open-source model SOLAR.

References

1. Auer, S., Bizer, C., Kobilarov, G., Lehmann, J., Cyganiak, R., Ives, Z.: Dbpedia: A nucleus for a web of open data. In: *Proceedings of the International Semantic Web Conference*. pp. 722–735. Springer (2007)
2. Cannavicchio, M., Barbosa, D., Merialdo, P.: Towards annotating relational data on the web with language models. In: *Proceedings of the 2018 World Wide Web Conference*. pp. 1307–1316 (2018)
3. Deng, X., Sun, H., Lees, A., Wu, Y., Yu, C.: TURL: Table understanding through representation learning. *Proceedings of the VLDB Endowment* **14**(3), 307–319 (Nov 2020)
4. Devlin, J., Chang, M.W., Lee, K., Toutanova, K.: BERT: Pre-training of deep bidirectional transformers for language understanding. In: *Proceedings of the 2019 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies, Volume 1*. pp. 4171–4186 (Jun 2019)
5. Feuer, B., Liu, Y., Hegde, C., Freire, J.: ArcheType: A Novel Framework for Open-Source Column Type Annotation using Large Language Models. *arXiv preprint arXiv:2310.18208* (2023)
6. Jiao, X., Yin, Y., Shang, L., Jiang, X., Chen, X., Li, L., et al.: TinyBERT: Distilling BERT for natural language understanding. In: *Findings of the Association for Computational Linguistics EMNLP 2020*. pp. 4163–4174 (Nov 2020)
7. Jiménez-Ruiz, E., Hassanzadeh, O., Efthymiou, V., Chen, J., Srinivas, K.: SemTab 2019: Resources to benchmark tabular data to knowledge graph matching systems. In: *Proceedings of the International Semantic Web Conference*. pp. 514–530 (2020)
8. Kayali, M., Lykov, A., Fountalis, I., Vasiloglou, N., Olteanu, D., Suciu, D.: CHO-RUS: Foundation Models for Unified Data Discovery and Exploration. *arXiv preprint arXiv:2306.09610* (2023)
9. Korini, K., Bizer, C.: Column type annotation using ChatGPT. In: *Joint proceedings of workshops at the 49th International Conference on Very Large Data Bases, CEUR-WS Vol. 3462*. pp. 1–12 (2023)
10. Korini, K., Peeters, R., Bizer, C.: SOTAB: The WDC schema.org table annotation benchmark. In: *Proceedings of the Semantic Web Challenge on Tabular Data to Knowledge Graph Matching (SemTab), CEUR-WS Vol. 3320*. pp. 14–19 (2022)
11. Li, P., He, Y., Yashar, D., Cui, W., Ge, S., Zhang, H., et al.: Table-GPT: Table-tuned GPT for Diverse Table Tasks. *arXiv preprint arXiv:2310.09263* (2023)
12. Limaye, G., Sarawagi, S., Chakrabarti, S.: Annotating and searching web tables using entities, types and relationships. *Proceedings of the VLDB Endowment* **3**(1-2), 1338–1347 (2010)
13. Liu, J., Shen, D., Zhang, Y., Dolan, W.B., Carin, L., Chen, W.: What makes good in-context examples for GPT-3? In: *Proceedings of Deep Learning Inside Out (DeeLIO 2022)*. pp. 100–114 (2022)
14. Liu, J., Chabot, Y., Troncy, R., Huynh, V.P., et al.: From tabular data to knowledge graphs: A survey of semantic table interpretation tasks and methods. *Journal of Web Semantics* **76**, 100761 (2023)
15. Liu, J., Troncy, R.: DAGOBAB: an end-to-end context-free tabular data semantic annotation system. In: *Proceedings of the Semantic Web Challenge on Tabular Data to Knowledge Graph Matching (SemTab), CEUR-WS Vol. 2553*. pp. 41–48 (2019)
16. Nguyen, P., Kertkeidkachorn, N., Ichise, R., Takeda, H.: Mtab: Matching tabular data to knowledge graph using probability models. In: *Proceedings of the Semantic*

- Web Challenge on Tabular Data to Knowledge Graph Matching (SemTab), CEUR-WS Vol. 2553. pp. 7–14 (2019)
17. Ouyang, L., Wu, J., Jiang, X., Almeida, D., Wainwright, C., Mishkin, P., et al.: Training language models to follow instructions with human feedback. *Advances in Neural Information Processing Systems* **35**, 27730–27744 (2022)
 18. Ritze, D., Lehmberg, O., Bizer, C.: Matching HTML tables to DBpedia. In: *Proceedings of the 5th International Conference on Web Intelligence, Mining and Semantics*. pp. 1–6 (2015)
 19. Suhara, Y., Li, J., Li, Y., Zhang, D., Demiralp, Ç., Chen, C., et al.: Annotating columns with pre-trained language models. In: *Proceedings of the 2022 International Conference on Management of Data*. pp. 1493–1503 (2022)
 20. Touvron, H., Martin, L., Stone, K., Albert, P., Almahairi, A., Babaei, Y., et al.: Llama 2: Open foundation and fine-tuned chat models. *arXiv preprint arXiv:2307.09288* (2023)
 21. Venetis, P., Halevy, A., Madhavan, J., Paşca, M., Shen, W., Wu, F., et al.: Recovering semantics of tables on the web. *Proceedings of the VLDB Endowment* **4**(9), 528–538 (2011)
 22. Ye, X., Iyer, S., Celikyilmaz, A., Stoyanov, V., Durrett, G., Pasunuru, R.: Complementary explanations for effective in-context learning. In: *Findings of the Association for Computational Linguistics: ACL 2023*. pp. 4469–4484 (2023)
 23. Zhang, T., Yue, X., Li, Y., Sun, H.: TableLlama: Towards Open Large Generalist Models for Tables. *arXiv preprint arXiv:2311.09206* (2023)
 24. Zhao, W.X., Zhou, K., Li, J., Tang, T., Wang, X., et al.: A Survey of Large Language Models. *arXiv preprint arXiv:2303.18223* (2023)